

# AI Presentation Coach

Using AI to make you a better public speaker

**An AWS-native event based agentic presentation coaching platform built with Kiro and AI-DLC practices.**

This is not just a demo, it is a useful pattern for taking agentic systems from idea to PoC to Production via governed software delivery.

Created and developed by Michael Geiser

Dark theme retained from the working deck



## AI Presentation Coach

Using AI to make you a better public speaker

### AI-Powered Presentation Coaching

Upload your presentation audio and receive detailed feedback across 7 dimensions: delivery, structure, executive presence, technical communication, audience engagement, pacing, and persuasion.

[Sign In to Get Started](#)

#### Presentation Coaching Report

User: Michael Geiser (mgeiser@mgeiser.net)  
Presentation Title: new test  
Description: test of Exec summary  
Presentation File: 1\_Introducing AgentCore.mp3  
Uploaded: June 22, 2026 21:42 ET

#### Executive Summary

This presentation is technically knowledgeable and genuinely engaging in tone, but it reads more like an enthusiastic product walkthrough than a confident, authoritative introduction — and that gap is what the 5.9 overall average reflects. That score signals a presenter with real substance who has not yet learned to carry that substance with full conviction. The content is credible, the energy is real, but the delivery patterns, structural looseness, and absence of persuasive architecture consistently blunt the impact of what is actually strong material.

The clearest strengths here are the technical depth and the instinct to make complex infrastructure concepts feel human. Delivery scores the highest of any dimension for good reason: the conversational warmth, inclusive second-person framing, and natural rhetorical questions in the memory section all create a sense of a knowledgeable guide rather than a lecturer reading slides. Technical Communication also scores well because the problem-first framing is genuinely effective — introducing each capability by naming the pain point it solves is sound instructional architecture. The honest acknowledgment of platform limitations and the proactive handling of the vendor-lock objection demonstrate a presenter who understands the audience's skepticism, and that kind of intellectual transparency builds real credibility with technical listeners.

The two highest-leverage improvements out across nearly every lower-scoring dimension simultaneously. First, Executive Presence and Persuasion both suffer from the same root cause: hedging and casual filler language — “kind of,” “sort of,” “super super easy,” “mishmash,” “if you ask me” — that systematically undercuts the authority this presenter has clearly earned. Replacing that language with precise, confident phrasing would lift scores across Delivery, Executive Presence, and Persuasion in a single pass. Second, Audience Engagement and Structure both suffer from the absence of a narrative spine. The presentation is currently a feature inventory with a loose framing device: transforming it into a problem-solution arc with a concrete protagonist — even a hypothetical developer struggling to get an agent into production — would give the audience something to follow emotionally, not just intellectually. Structure also needs a clean verbal agenda at the opening and a proper 30-to-45-second closing recap, rather than the current 12-second pivot to the demo. For the next recording, focus on three things: open with one bold, evidence-anchored statement of impact instead of “hey everyone,” replace every instance of “kind of” and “sort of” with either a precise word or a clean pause, and close with a 40-second summary that names the three capabilities the audience should remember and tells them exactly what they are about to do in the hands-on session and why it matters.

**AWS-native  
serverless +  
agentic**

**Built for  
zero idle cost**

## EXECUTIVE LENS

# The CTO readout

The AI Presentation Coach (“PresCoach”) is a working AWS-native product pattern, not a lab toy. It accepts audio, prepares it deterministically, then hands it to independent Agentic AI evaluators that produce a coaching report.

**7**

AI coaching dimensions

**20+**AWS native  
services used**726**

Tests through all workflows

**\$33.69**Spend for environment  
and AI tools used

(\$6.59 AWS + \$27.10 Kiro)

**Architecture decision  
worth a CTO review****Where to harden  
before real users****How to scale  
without cost traps****What AI-DLC  
actually changed**

## WHAT IT DOES

# Using the product simple, the system behind it is not

**Presentation Coach** Administrator Upload Submissions Logout

### Upload Presentation

Presentation File \* Required

Choose File 2\_Andy Jassy AgentCore.mp3

2\_Andy Jassy AgentCore.mp3 (1.97 MB)

Presentation Title \* Required

test for #2 11 / 200

Description (optional)

Describe your presentation

**Presentation Coach** Administrator Upload Submissions Logout

### My Submissions

|                                                                       |               |
|-----------------------------------------------------------------------|---------------|
| test for #2                                                           | Pending       |
| File: 2_Andy Jassy AgentCore.mp3<br>Uploaded: Jun 25, 2026, 10:49 AM  | Delete        |
| Test for #1                                                           | Processing... |
| File: 1_Introducing AgentCore.mp3<br>Uploaded: Jun 25, 2026, 10:47 AM | Evaluating    |
|                                                                       | Delete        |
| test of Pres #4                                                       | Completed     |

## User-facing flow



The deliberate split: deterministic processing does the mechanical work, while the probabilistic agentic layer handles judgment. That separation is what makes the design easier to govern.


Let's Deep Dive into the App UI Frontend & Backend, processing, and Agentic layers and then review the AI-DLC behind it.

# The Web App

# Responsive Web App




Web



Mobile

# Cognito Based User IAM



## Groups:

- Users
- Administrators

## Other Cognito Functions implemented

- User self signup
- Password reset
- Password recovery
- MFA not implemented in MVP

# Audio File Upload

 **Presentation Coach**

[Administrator](#) [Upload](#) [Submissions](#) [Logout](#)

## Upload Presentation

Presentation File \* Required

Choose File

1\_Introducing AgentCore.mp3

1\_Introducing AgentCore.mp3 (2.45 MB)

Presentation Title \* Required

test run

9 / 200

Description (optional)

Describe your presentation

0 / 2000

Upload

 © 2026 Geislersoft Presentation Coaching Platform

## Features

- File picker
- Title
- Description
- Video File upload on roadmap

# Processing

|                                                                                                             |  |                               |                        |                             |                        |
|-------------------------------------------------------------------------------------------------------------|--|-------------------------------|------------------------|-----------------------------|------------------------|
|  <b>Presentation Coach</b> |  | <a href="#">Administrator</a> | <a href="#">Upload</a> | <a href="#">Submissions</a> | <a href="#">Logout</a> |
| <b>My Submissions</b>                                                                                       |  |                               |                        |                             |                        |
| <b>test 3</b>                                                                                               |  | Pending                       |                        |                             |                        |
| File: 3_Agentcore-Introduction.mp3<br>AgentCore Introduct from AWS<br>Uploaded: Jun 20, 2026, 02:37 PM      |  | Delete                        |                        |                             |                        |
| <b>test of another presentation</b>                                                                         |  | Completed                     |                        |                             |                        |
| File: 2_Andy lassy AgentCore.mp3<br>Uploaded: Jun 20, 2026, 11:12 AM                                        |  | Delete                        |                        |                             |                        |
| Download Report                                                                                             |  |                               |                        |                             |                        |
| <b>test repeat</b>                                                                                          |  | Completed                     |                        |                             |                        |
| File: 1_Introducing AgentCore.mp3<br>Uploaded: Jun 20, 2026, 09:07 AM                                       |  | Delete                        |                        |                             |                        |
| Download Report                                                                                             |  |                               |                        |                             |                        |
| <b>test</b>                                                                                                 |  | Failed                        |                        |                             |                        |
| File: 1_Introducing AgentCore.mp3<br>Uploaded: Jun 19, 2026, 02:52 PM                                       |  | Delete                        |                        |                             |                        |

## Features

- Multiple steps
- Realtime User status updates



# Statuses

## Start of Presentation Coach Process

|                                                                                                                                                        |               |            |
|--------------------------------------------------------------------------------------------------------------------------------------------------------|---------------|------------|
| <div>Report test</div> <div>File: 1_Introducing AgentCore.mp3</div> <div>Verifying the report output</div> <div>Uploaded: Jun 22, 2026, 02:05 PM</div> | Pending       | Delete     |
| <div>Report test</div> <div>File: 1_Introducing AgentCore.mp3</div> <div>Verifying the report output</div> <div>Uploaded: Jun 22, 2026, 02:05 PM</div> | Processing... | Delete     |
| <div>Report test</div> <div>File: 1_Introducing AgentCore.mp3</div> <div>Verifying the report output</div> <div>Uploaded: Jun 22, 2026, 02:05 PM</div> | Waiting       | Delete     |
| <div>Report test</div> <div>File: 1_Introducing AgentCore.mp3</div> <div>Verifying the report output</div> <div>Uploaded: Jun 22, 2026, 02:05 PM</div> | Processing... | Evaluating |

**Pending:** User uploaded unprocessed audio file

**Processing:** Transcription and prep

**Waiting:** Queued for Agentic Evaluation

**Evaluation:** Agentic analysis in progress

# Statuses

|                                                                                                                            |               |                   |        |
|----------------------------------------------------------------------------------------------------------------------------|---------------|-------------------|--------|
| <b>new test</b><br>File: 1_Introducing AgentCore.mp3<br>test of Exec summary<br>Uploaded: Jun 22, 2026, 09:42 PM           | Processing... | Generating Report | Delete |
| <b>Report test</b><br>File: 1_Introducing AgentCore.mp3<br>Verifying the report output<br>Uploaded: Jun 22, 2026, 02:05 PM |               | Completed         | Delete |

## End of “Happy Path” Presentation Coach Process

|                                                                                      |  |        |        |
|--------------------------------------------------------------------------------------|--|--------|--------|
| <b>test</b><br>File: 1_Introducing AgentCore.mp3<br>Uploaded: Jun 19, 2026, 02:52 PM |  | Failed | Delete |
|--------------------------------------------------------------------------------------|--|--------|--------|

**Generating Report:**  
Final report is being created

**Completed:** Report is ready for download

**Failed:** An error has halted processing

# Report is Ready to Read

|                                                                                                             |  |                        |        |             |        |
|-------------------------------------------------------------------------------------------------------------|--|------------------------|--------|-------------|--------|
|  <b>Presentation Coach</b> |  | <b>Administrator</b>   | Upload | Submissions | Logout |
| <b>My Submissions</b>                                                                                       |  |                        |        |             |        |
| <b>test 3</b>                                                                                               |  | Completed              |        |             |        |
| Files: 3_AgentCore-Introduction.mp3<br>AgentCore Introduce from AWS<br>Uploaded: Jun 20, 2026, 02:37 PM     |  |                        |        |             |        |
| <a href="#">Download Report</a>                                                                             |  | <a href="#">Delete</a> |        |             |        |
| <b>test of another presentation</b>                                                                         |  | Completed              |        |             |        |
| Files: 2_Andy.Jassy.AgentCore.mp3<br>Uploaded: Jun 20, 2026, 11:12 AM                                       |  |                        |        |             |        |
| <a href="#">Download Report</a>                                                                             |  | <a href="#">Delete</a> |        |             |        |
| <b>test repeat</b>                                                                                          |  | Completed              |        |             |        |
| Files: 1_Introducing AgentCore.mp3<br>Uploaded: Jun 20, 2026, 09:07 AM                                      |  |                        |        |             |        |
| <a href="#">Download Report</a>                                                                             |  | <a href="#">Delete</a> |        |             |        |
| <b>test</b>                                                                                                 |  | Failed                 |        |             |        |
| Files: 1_Introducing AgentCore.mp3<br>Uploaded: Jun 19, 2026, 02:52 PM                                      |  |                        |        |             |        |
|                                                                                                             |  | <a href="#">Delete</a> |        |             |        |

## Features

- Current status of analysis
- Data entered by User displayed
- Download report link
- Delete link

# The Presentation Coaching Report

## Presentation Coaching Report

**User:** Michael Geiser (mgeiser@mgeiser.net)  
**Presentation Title:** new test  
**Description:** test of Exec summary  
**Presentation File:** 1\_Introducing AgentCore.mp3  
**Uploaded:** June 22, 2026 21:42 ET

### Executive Summary

This presentation is technically knowledgeable and genuinely engaging in tone, but it reads more like an enthusiastic product walkthrough than a confident, authoritative introduction — and that gap is what the 5.9 overall average reflects. That score signals a presenter with real substance who has not yet learned to carry that substance with full conviction. The content is credible, the energy is real, but the delivery patterns, structural looseness, and absence of persuasive architecture consistently blunt the impact of what is actually strong material.

The clearest strengths here are the technical depth and the instinct to make complex infrastructure concepts feel human. Delivery scores the highest of any dimension for good reason: the conversational warmth, inclusive second-person framing, and natural rhetorical questions in the memory section all create a sense of a knowledgeable guide rather than a lecturer reading slides. Technical Communication also scores well because the problem-first framing is genuinely effective — introducing each capability by naming the pain point it solves is sound instructional architecture. The honest acknowledgment of platform limitations and the proactive handling of the vendor-lock objection demonstrate a presenter who understands the audience's skepticism, and that kind of intellectual transparency builds real credibility with technical listeners.

The two highest-leverage improvements cut across nearly every lower-scoring dimension simultaneously. First, Executive Presence and Persuasion both suffer from the same root cause: hedging and casual filler language — "kind of," "sort of," "super super easy," "mishmash," "if you ask me" — that systematically undercuts the authority this presenter has clearly earned. Replacing that language with precise, confident phrasing would lift scores across Delivery, Executive Presence, and Persuasion in a single pass. Second, Audience Engagement and Structure both suffer from the absence of a narrative spine. The presentation is currently a feature inventory with a loose framing device; transforming it into a problem-solution arc with a concrete protagonist — even a hypothetical developer struggling to get an agent into production — would give the audience something to follow emotionally, not just intellectually. Structure also needs a clean verbal agenda at the opening and a proper 30-to-45-second closing recap, rather than the current 12-second pivot to the demo. For the next recording, focus on three things: open with one bold, evidence-anchored statement of impact instead of "Hey everyone," replace every instance of "kind of" and "sort of" with either a precise word or a clean pause, and close with a 40-second summary that names the three capabilities the audience should remember and tells them exactly what they are about to do in the hands-on session and why it matters.

### Per-Dimension Detailed Feedback

#### Delivery (Score: 6.8/10.0)

##### Findings:


- [MEDIUM] Filler Words: Two instances of 'uh' were detected in the transcript (at approximately 159s and 241s). Additionally, conversational hedges like 'you know' (at ~38s), 'kind of' (used multiple times), and 'sort of' appear several times throughout the presentation, reducing overall polish.  
→ Suggestion: Practice pausing silently instead of using filler words. Record yourself and flag each 'uh,' 'you know,' and 'kind of' — replace them with a deliberate 1-2 second pause to maintain authority and clarity.
- [MEDIUM] Pace: The presentation spans approximately 307 seconds (just over 5 minutes) and covers a dense amount of technical content. Analysis of word timing shows a largely consistent and moderately fast pace with few natural breaks between major topic transitions. Some segments, like the capabilities overview (179s–267s), rush through multiple distinct features without sufficient pacing variation.  
→ Suggestion: Intentionally slow down when introducing a new concept or feature (e.g., identity, memory, gateways). Aim for a 10–20% pace reduction on key points to allow the audience to absorb information before moving on.
- [MEDIUM] Pauses: There are some natural pauses at sentence boundaries (e.g., after 'It's really cool.' at ~9.7s, and after 'Big one is memory.' at ~221s), but strategic mid-section pauses for emphasis are largely absent. Transitions between major capability sections (identity → authentication → tools → memory → gateways) lack deliberate pause moments that signal a shift in topic.  
→ Suggestion: Insert 2–3 second intentional pauses before introducing each major capability (e.g., before 'Big one is memory'). These signpost pauses help the audience mentally prepare to receive new information and add dramatic effect to key points.
- [MEDIUM] Vocal Variety: The delivery is conversational and relatively smooth, but the transcript shows a somewhat uniform rhythm throughout. High-value phrases such as 'production-level solutions,' 'real business value,' and 'Big one is memory' could benefit from stronger tonal emphasis. The use of informal language ('mishmash of stuff,' 'fuss around,' 'nitty gritty') adds personality but may dilute authority when not paired with vocal emphasis.  
→ Suggestion: Identify 4–5 key phrases per section and rehearse delivering them with elevated pitch or increased volume. For example, slow down and drop pitch on 'that's not gonna cut it at scale in production' to underscore importance.
- [LOW] Energy: The presenter demonstrates genuine enthusiasm, particularly evident in phrases like 'It's really cool,' 'that's pretty cool if you ask me,' and 'super, super easy to deal with.' The casual, conversational tone conveys authenticity and approachability throughout the 307-second presentation.  
→ Suggestion: Channel current enthusiasm more strategically — save the highest energy moments for the most critical benefits (e.g., when explaining memory management or the framework-agnostic nature of

## Features

Detailed feedback across 7 dimensions:

- Delivery
- Structure
- Executive presence
- Technical communication
- Audience engagement
- Pacing
- Persuasion

# CICS Pipelines




The screenshot shows the AWS CodePipeline console interface. On the left is a navigation sidebar with 'Developer Tools' and 'CodePipeline'. The main area is titled 'Pipelines' and contains a table of pipeline details.

| Name                                         | Latest execution status | Latest source revisions                                                                                | Latest execution started | Most recent executions |
|----------------------------------------------|-------------------------|--------------------------------------------------------------------------------------------------------|--------------------------|------------------------|
| prescoach-dev-kiro-eval-workflow-full-deploy | Succeeded               | GitHub - e187925d: fix: corecting erros on eval work unit tests                                        | 2 days ago               | View details           |
| prescoach-dev-kiro-eval-workflow-deploy      | Succeeded               | GitHub - 5743dcde: exec summary changes                                                                | 2 days ago               | View details           |
| prescoach-dev-kiro-eval-workflow-test        | Succeeded               | GitHub - 358e4f6b: fix: use cross-region inference profile us.anthropic.claude-sonnet-4-6 (required... | 7 days ago               | View details           |
| prescoach-dev-kiro-prep-workflow-full-deploy | Succeeded               | GitHub - 30c80d86: fix: use existing uploads bucket for vector storage instead of nonexistent presc... | 9 days ago               | View details           |
| prescoach-dev-kiro-prep-workflow-test        | Succeeded               | GitHub - 358e4f6b: fix: use cross-region inference profile us.anthropic.claude-sonnet-4-6 (required... | 7 days ago               | View details           |
| prescoach-dev-kiro-prep-workflow-deploy      | Succeeded               | GitHub - e0380c9a: fix: status and processing layout                                                   | 4 days ago               | View details           |
| prescoach-dev-kiro-frontend                  | Succeeded               | GitHub - 7840c001: html: UI tweaks                                                                     | 4 days ago               | View details           |
| prescoach-dev-kiro-backend                   | Succeeded               | GitHub - e0380c9a: fix: status and processing layout                                                   | 4 days ago               | View details           |
| prescoach-dev-kiro-full-deploy               | Succeeded               | GitHub - f53e6401: feat: add report download link — generate presigned S3 URL for completed submiss... | 7 days ago               | View details           |

CICD Pipelines allow deploys at logical unit level:


- Frontend/Backend
- Audio/Video Processing
- Agentic AI Evaluation



Each Pipeline can deploy

- Tests only ("test")
- Code only ("deploy")
- Tests & code ("full-deploy")


# Development Costs




## Features

Total costs of the development

- AWS Account: \$6.59
- Kiro AI-DLC: 1355 Credits (\$27.10)



# Github Repo Activity




<https://github.com/michaelgeiser/kiro-agentcore-app>



# The Architecture





# Web App SPA and Backend



# Audio File Prep




# Evaluation with Agentic AI and Report Creation



## DELIVERY MODEL

# Build approach: AI-DLC turned the idea into deployable work units


The useful lesson is not that code was generated quickly. The useful lesson is that the work was decomposed into independently testable “bolts” with clear contracts.



Retrospective describes four work units, bolt-based interfaces, CI/CD pipelines, and zero-idle service selection.

## RUNTIME FLOW

# End-to-end processing is event-driven and observable




The handoff contract matters more than the tool choice. The upload service does not know how evaluation works. The evaluator does not care how the file arrived. That keeps each boundary testable and replaceable.

**State transitions**  
**Pending** → **Processing** → **Completed**  
or **Failed**

**Failure path**  
**retries** → **HandleFailure** → **SNS + DLQ**

# Architecture, viewed like a CTO would review it



## Why this is a good MVP architecture

- Files never pass through Lambda
- Async processing gives back-pressure
- Feature flags allow runtime tuning
- Least-privilege IAM can be reviewed per work unit

## What I would challenge before production

- Model quotas and throttling behavior
- Regional failure stance and user-pool strategy
- PII retention and report lifecycle policy
- Evaluation result explainability and traceability


Architecture summarizes services described in project overview, flow, and retrospective.

# The agentic layer is composable by design

This is the part that justifies an agentic approach. The value comes from multiple independent lenses reviewing the same asset, not from one big prompt pretending to be a committee.

Agent definitions live in a JSON manifest and can be enabled or disabled without code changes.

The report generator assembles structured results into a 10-18 page PDF coaching report.



## COST AND SCALE

# The economics are exactly why serverless fits

For spiky, early-stage usage, the cost goal is not “cheapest at theoretical peak.” It is “do not pay for infrastructure while product-market fit is still loading” but will scale economically as usage takes off.

**Demo / PoC**

10/month

 **\$0.55****Regular use**

500/month

 **\$28****Conference spike**

10,000/week

 **\$531****Quiet after spike**

following month

 **\$2.35**

**Bedrock dominates  
the bill at scale**

**Zero idle  
for Lambda, API Gateway,  
Step Functions, SQS, SNS,  
Bedrock**

**S3 vector store  
keeps MVP cost near zero**

Cost scenarios and zero-idle principle come from project retrospective cost analysis.



## COST DECISION

# S3 as the first vector store is the right kind of cheap

This is one of those choices where CTO discipline matters. For the MVP, S3 is not the “best vector database.” It is the best way to avoid paying for one before the workload proves it deserves one; it is **fit to purpose** engineering.

## S3 JSON embeddings

**Idle cost**

\$0 storage only

**Read pattern**

Per-submission retrieval

**Latency**

~50-100 ms per file

**Move when**

MVP and variable load

## OpenSearch Serverless / pgvector

~\$175/month minimum

Similarity search across corpora

~10 ms query path


&gt;50K presentations/month or cross-submission search

Design win: vector-store-type and vector-store-endpoint are SSM parameters, so the upgrade path is operationally clean.

## RESILIENCE

## Operational resilience is solid for MVP and launch, but the regional stance is a decision

Current state is single-region in us-east-1 and multi-AZ compatible by design and construction. Single Region is acceptable for an MVP, but production requires an explicit answer to how much regional failure the business will tolerate. Multi-region is cost effective and does not require re-architecture work.



Low-friction upgrades to multi-region resilience: DynamoDB Global Tables, S3 CRR, second regional API Gateway, replicated Step Functions and Lambda stacks, CloudFront origin routing.

## RISK REGISTER

# The production-hardening backlog is manageable, not mysterious

The system is deployed and the end-to-end pipeline runs. The remaining work is application-level correctness plus the usual enterprise-grade hardening. Good news: no smoke machine required.

**Known app fixes**

Use Web Sockets instead of polling for Submissions page while a presentation is processing.

**Security**

Add additional guardrails, add admin approvals for new users, enhance WAF rules, define data retention policies.

**Reliability**

Refine throttling policy by user and entire system, ECS Spot interruption handling, DLQ message runbooks, retry budgets, test idempotency fully.

**Observability**

Add Correlation Id and additional logging to trace a submission across Cognito, API, S3, Step Functions, queues, ECS, Bedrock, and report generation.

**FinOps**

Refine budgets, Bedrock spend alarms, lifecycle policies, model-choice controls, chargeback by user or tenant.


Evaluate items and prioritize:  
In this case the Security and Reliability issues have the highest priority

Known application issues come from project overview and source deck. Hardening items are CTO recommendations derived from the architecture.

## WHERE I WOULD SPEND TIME

# CTO decision radar

The architecture is good enough to advance. The discussion should move from “can this work?” to “what operating posture do we want?”



This slide distills CTO decisions from the supplied architecture, flow, cost, and current-state documents.

## NEXT MOVES

# A practical 15/30/60 day path

Do not boil the architecture ocean. Fix the known seams, prove the evaluator quality, then harden only the parts that have a user or revenue reason to exist.

**15 days****Stabilize MVP**

Refine throttling policy  
ECS Spot interruption handling,  
Add end-to-end trace ID  
Budget alarms on Bedrock usage

**30 days****Prove quality**

Golden set for coaching quality  
Prompt and model versioning  
Agent output schema enforcement  
Report scoring calibration

**60 days****Prepare production posture**

Retention and PII controls  
RTO/RPO decision  
Quotas and load test  
Tenant or user-level cost attribution

**Exit criteria**

A CTO should be able to trace one submission from upload to report, explain one agent's score, and predict the cost of the next 1,000 runs.

CLOSE

# The CTO-to-CTO takeaway

**PresCoach is a strong pattern because it combines three things that usually get separated: a useful product experience, an AWS-native cost model, and an AI-DLC delivery method that keeps architecture boundaries visible.**

## What is working

Clear work-unit boundaries, async flow, managed services, zero idle cost, composable agents.

## What to fix

Retrieval, parser handling, status page behavior, traceability, hardening of AI and data controls.

## What to decide

Regional posture, vector store timing, quality gates, model spend controls, production-readiness threshold.

**Bottom line: advance it but make the next phase *boring* in the best possible enterprise way. Boring is where systems go to become dependable.**


# Appendix Material

# Architecture Flow Diagrams




# Upload and Storage Service — Architecture Diagram

prescoach-dev-kiro




# Preparation Workflow — Architecture Diagram

prescoach-dev-kiro | Step Functions Standard Workflow | 10 Lambda Functions



Agentic Evaluation — LOCAL MODE Architecture

ECS Fargate Spot | Strands SDK Local Execution | Direct Bedrock InvokeModel



**LOCAL MODE — Key Characteristics**

- All agents run in-process inside a single ECS Fargate Spot container (no AgentCore deployment required)
- Strands SDK creates Agent objects locally — each evaluation tool calls bedrock:InvokeModel directly
- Session state is in-memory only (lost when task exits) — no persistent memory across task launches
- Cost-optimized: Fargate Spot (70% discount), 30-min idle timeout, concurrent processing of 5 submissions
- Cold start: ~45-90s (container pull + startup). Kept warm for 30 min after last message.

**Cost Estimate**

Per evaluation task (5 min): ~\$0.0007

Monthly @ 5/day: ~\$0.10

Monthly @ 100/day: ~\$2.10

(Bedrock model costs are additional)

# Agentic Evaluation — AGENTCORE MODE Architecture

Amazon Bedrock AgentCore Runtime | Managed Agents | Session Memory | Auto-Scaling

